

```
In [ ]: # Problem :
# Part 1: Intelligent Lead Scoring & CLM Prediction (Python/Scala)
# Objective: Predict Lead conversion and manage the full client lifecycle using AI.

# Tasks:
# - Build a supervised ML model to score Leads based on various engagement metrics.
# - Build a second model to predict customer churn or next best action (NBA).
# -Train, evaluate, and deploy via REST API (Flask/FastAPI).

# Bonus:
# - Output lifecycle stage prediction.
# - Visualize data.
# - Use Spark/Scala if needed
```

```
In [1]: #Import necessary Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# Create Dataframe
My_Data= {
    'lead_source': ['Website', 'Email', 'Phone', 'Website', 'Email', 'Phone'] * 20,
    'engagement_score': [75, 45, 60, 80, 30, 55] * 20,
    'pages_visited': [5, 2, 3, 6, 1, 4] * 20,
    'time_spent': [300, 120, 200, 400, 90, 250] * 20,
    'converted': [1, 0, 1, 1, 0, 0] * 20,
    'churned': [0, 1, 0, 0, 1, 1] * 20
}

Data = pd.DataFrame(My_Data)

# Convert the datatype of Lead_source
Data['lead_source'] = Data['lead_source'].astype('category').cat.codes

# Create Target column for Lead_Score using existing column in dataset
X_lead = Data[['lead_source', 'engagement_score', 'pages_visited', 'time_spent']]
y_lead = Data['converted']

# create Target column for churn prediction using existing columns in dataset
X_churn = Data[['lead_source', 'engagement_score', 'pages_visited', 'time_spent']]
y_churn = Data['churned']

# Split the datasets
X_train_lead, X_test_lead, y_train_lead, y_test_lead = train_test_split(X_lead, y_lead)
X_train_churn, X_test_churn, y_train_churn, y_test_churn = train_test_split(X_churn, y_churn)

# Train Lead scoring model using Random Forest
lead_model = RandomForestClassifier(random_state=42)
lead_model.fit(X_train_lead, y_train_lead)

# Train churn prediction model
```

```
churn_model = RandomForestClassifier(random_state=42)
churn_model.fit(X_train_churn, y_train_churn)

# Evaluate models
lead_preds = lead_model.predict(X_test_lead)
churn_preds = churn_model.predict(X_test_churn)

lead_report = classification_report(y_test_lead, lead_preds, output_dict=True)
churn_report = classification_report(y_test_churn, churn_preds, output_dict=True)

lead_report, churn_report
```

```
Out[1]: ({'0': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 16.0},
          '1': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 8.0},
          'accuracy': 1.0,
          'macro avg': {'precision': 1.0,
                        'recall': 1.0,
                        'f1-score': 1.0,
                        'support': 24.0},
          'weighted avg': {'precision': 1.0,
                           'recall': 1.0,
                           'f1-score': 1.0,
                           'support': 24.0}},
         {'0': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 8.0},
          '1': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 16.0},
          'accuracy': 1.0,
          'macro avg': {'precision': 1.0,
                        'recall': 1.0,
                        'f1-score': 1.0,
                        'support': 24.0},
          'weighted avg': {'precision': 1.0,
                           'recall': 1.0,
                           'f1-score': 1.0,
                           'support': 24.0}}})
```

In [3]: Data

Out[3]:

	lead_source	engagement_score	pages_visited	time_spent	converted	churned
0	2	75	5	300	1	0
1	0	45	2	120	0	1
2	1	60	3	200	1	0
3	2	80	6	400	1	0
4	0	30	1	90	0	1
...
115	0	45	2	120	0	1
116	1	60	3	200	1	0
117	2	80	6	400	1	0
118	0	30	1	90	0	1
119	1	55	4	250	0	1

120 rows × 6 columns

In [4]: `pip install flask scikit-learn pandas matplotlib seaborn`

Requirement already satisfied: flask in c:\users\dell\anaconda3\lib\site-packages (3.0.3)

Requirement already satisfied: scikit-learn in c:\users\dell\anaconda3\lib\site-packages (1.5.1)

Requirement already satisfied: pandas in c:\users\dell\anaconda3\lib\site-packages (2.2.2)

Requirement already satisfied: matplotlib in c:\users\dell\anaconda3\lib\site-packages (3.9.2)

Requirement already satisfied: seaborn in c:\users\dell\anaconda3\lib\site-packages (0.13.2)

Requirement already satisfied: Werkzeug>=3.0.0 in c:\users\dell\anaconda3\lib\site-packages (from flask) (3.0.3)

Requirement already satisfied: Jinja2>=3.1.2 in c:\users\dell\anaconda3\lib\site-packages (from flask) (3.1.4)

Requirement already satisfied: itsdangerous>=2.1.2 in c:\users\dell\anaconda3\lib\site-packages (from flask) (2.2.0)

Requirement already satisfied: click>=8.1.3 in c:\users\dell\anaconda3\lib\site-packages (from flask) (8.1.7)

Requirement already satisfied: blinker>=1.6.2 in c:\users\dell\anaconda3\lib\site-packages (from flask) (1.6.2)

Requirement already satisfied: numpy>=1.19.5 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn) (1.26.4)

Requirement already satisfied: scipy>=1.6.0 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn) (1.13.1)

Requirement already satisfied: joblib>=1.2.0 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\dell\anaconda3\lib\site-packages (from scikit-learn) (3.5.0)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\dell\anaconda3\lib\site-packages (from pandas) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in c:\users\dell\anaconda3\lib\site-packages (from pandas) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in c:\users\dell\anaconda3\lib\site-packages (from pandas) (2023.3)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (1.2.0)

Requirement already satisfied: cycler>=0.10 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (4.51.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (1.4.4)

Requirement already satisfied: packaging>=20.0 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (24.1)

Requirement already satisfied: pillow>=8 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (10.4.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\dell\anaconda3\lib\site-packages (from matplotlib) (3.1.2)

Requirement already satisfied: colorama in c:\users\dell\anaconda3\lib\site-packages (from click>=8.1.3->flask) (0.4.6)

Requirement already satisfied: MarkupSafe>=2.0 in c:\users\dell\anaconda3\lib\site-packages (from Jinja2>=3.1.2->flask) (2.1.3)

Requirement already satisfied: six>=1.5 in c:\users\dell\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)

Note: you may need to restart the kernel to use updated packages.

```

In [5]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.ensemble import RandomForestClassifier
        from flask import Flask, request, jsonify
        import joblib

        # Dummy data generation
        def generate_data():
            data = {
                'lead_source': ['Website', 'Email', 'Phone'] * 40,
                'engagement_score': [75, 45, 60] * 40,
                'pages_visited': [5, 2, 3] * 40,
                'time_spent': [300, 120, 200] * 40,
                'converted': [1, 0, 1] * 40,
                'churned': [0, 1, 0] * 40
            }
            df = pd.DataFrame(data)
            df['lead_source'] = df['lead_source'].astype('category').cat.codes
            return df

        df = generate_data()

        # Train Lead Scoring Model
        X_lead = df[['lead_source', 'engagement_score', 'pages_visited', 'time_spent']]
        y_lead = df['converted']
        X_train_lead, _, y_train_lead, _ = train_test_split(X_lead, y_lead, test_size=0.2)
        lead_model = RandomForestClassifier().fit(X_train_lead, y_train_lead)
        joblib.dump(lead_model, 'lead_model.pkl')

        # Train Churn Prediction Model
        y_churn = df['churned']
        X_train_churn, _, y_train_churn, _ = train_test_split(X_lead, y_churn, test_size=0.2)
        churn_model = RandomForestClassifier().fit(X_train_churn, y_train_churn)
        joblib.dump(churn_model, 'churn_model.pkl')

```

Out[5]: ['churn_model.pkl']

```

In [7]: from flask import Flask, request, jsonify
        import joblib

        app = Flask(__name__)
        lead_model = joblib.load('lead_model.pkl')
        churn_model = joblib.load('churn_model.pkl')

        @app.route('/predict-lead', methods=['POST'])
        def predict_lead():
            data = request.json
            pred = lead_model.predict([list(data.values())])[0]
            return jsonify({'lead_conversion': int(pred)})

        @app.route('/predict-churn', methods=['POST'])
        def predict_churn():
            data = request.json
            pred = churn_model.predict([list(data.values())])[0]
            return jsonify({'churn_prediction': int(pred)})

```

```
if __name__ == '__main__':
    app.run(debug=True)
```

```
* Serving Flask app '__main__'
* Debug mode: on
```

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

```
* Running on http://127.0.0.1:5000
```

Press CTRL+C to quit

```
* Restarting with watchdog (windowsapi)
```

An exception has occurred, use %tb to see the full traceback.

SystemExit: 1

```
In [8]: import seaborn as sns
import matplotlib.pyplot as plt

# Correlation heatmap
sns.heatmap(df.corr(), annot=True)
plt.title("Feature Correlation Heatmap")
plt.show()
```

